



ARTIFICIAL NEURAL NETWORKS - WEEK 9

Long Short-Term Memory (LSTM)

Dr. Aamir Akbar

Director of both [AWKUM AI Lab](#) and [AWKUM Robotics](#), [Final Year Projects \(FYPs\)](#) coordinator, and lecturer at the department of Computer Science
Abdul Wali Khan University, Mardan (AWKUM)

CONTENTS

1. Introduction
2. LSTM Cell and Gates Logic
3. Case Study: Ethereum Price Prediction
4. Resources

LONG SHORT-TERM MEMORY (LSTM)


What is LSTM?

LSTM is a type of RNN architecture. It's designed to overcome the limitations of traditional RNNs in capturing and remembering long-term dependencies in sequential data.


In LSTM networks, each cell has a more complex structure compared to a basic RNN unit. It contains a memory cell that can maintain information over time, as well as various gates that regulate the flow of information into and out of the cell. These gates include the input gate, forget gate, and output gate.

1. **Input gate:** it determines how much new information is allowed into the memory cell.
2. **Forget gate:** it decides which information to discard from the memory cell.
3. **Output gate:** it controls the extent to which the information in the memory cell is used to compute the output of the LSTM unit.

LSTM ARCHITECTURE



LSTM ARCHITECTURE



sigmoid



tanh

pointwise
multiplicationpointwise
additionvector
concatenation


FORGOT GATE

Why Forget Gate is required?

1. **Long-Term Dependencies:** In sequences of data, there are often long-term dependencies where certain information becomes less relevant as the sequence progresses. The forget gate allows the LSTM to selectively discard information from the past that is no longer relevant for the current context.
2. **Addressing Vanishing Gradient:** Traditional RNNs often suffer from the vanishing gradient problem, where gradients diminish exponentially as they propagate back through time during training. By selectively forgetting irrelevant information through the forget gate, LSTMs mitigate the vanishing gradient problem by allowing the network to focus on relevant information, thereby preserving the gradients that contribute to learning meaningful patterns in the data.
3. **Dynamic Memory Management::** The forget gate enables LSTMs to dynamically manage their memory by updating and maintaining only the most relevant information in the memory cell while discarding less important information.

FORGOT GATE

Customers Review 2,491



Thanos

September 2018

Verified Purchase


Amazing! This box of cereal gave me a perfectly balanced breakfast, as all things should be. I only ate half of it but will definitely be buying again!



A Box of Cereal
\$3.99

FORGOT GATE

Customers Review 2,491



Thanos

September 2018

Verified Purchase

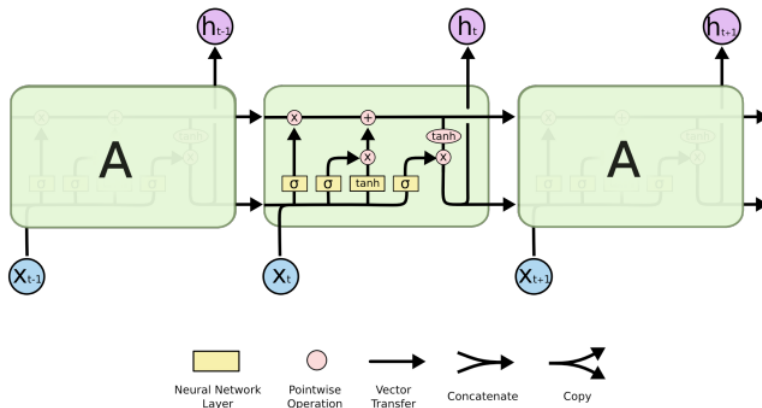
Amazing! This box of cereal gave me a perfectly balanced breakfast, as all things should be. I only ate half of it but will definitely be buying again!



A Box of Cereal
\$3.99

LSTM LAYERS



An LSTM is composed of four different layers.



LSTM CELL STATE OR MEMORY

The cell state is kind of like a conveyor belt.

The LSTM does have the ability to remove or add information to the cell state, carefully regulated by structures called gates.




Gates are composed out of a sigmoid neural net layer and a pointwise multiplication operation.

The sigmoid layer outputs numbers between **zero and one**, describing how much of each component should be let through. A value of **zero means let nothing through**, while a value of **one means let everything through**.

LSTM FORGET GATE

Should we continue to remember this bit of information or not?

$f_t = 1$ represents **completely keep this** while a 0 represents **completely get rid of this**.




$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Example: we want to predict the next word based on all the previous ones. In such a problem, the cell state might include the gender of the present subject, so that the correct pronouns can be used. When we see a new subject, we want to forget the gender of the old subject.

LSTM INPUT GATE

Should we update this bit of information or not?

If so, with what?




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Example: in the example of our language model, we want to add the gender of the new subject to the cell state, to replace the old one we're forgetting.

LSTM MEMORY UPDATE

- Forget that + memorize this



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$


Example: in the case of the language model, this is where we will actually drop the information about the old subject's gender and add the new one.

LSTM OUTPUT GATE

What shall we output given the context?

The output is a function of the cell state (C_t). A **tanh** pushes the cell state values towards $+1$ or -1 .

A **Perceptron** selects what to output based on the current input x_t and the previous output h_{t-1} , by multiplying its output o_t and $\tanh(C_t)$.




$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Example: in the case of the language model, given that we just saw a new subject, this cell might decide how to conjugate a verb, based on its knowledge of the subject, whether it is plural or singular.

THE LSTM BUILDING BLOCKS



- There are three gates (i, f, o) controlled by NNs (e.g., Perceptrons) weighing the inputs and the recurrent outputs.
- Given an input IN , we compute a new state \tilde{C} that can:
 1. be ignored ($i = 0$)
 2. replace the previous state ($f = 1$ and $i = 1$).
 3. be used to compute a new state $C + \tilde{C}$.
- Given a state C , the network outputs OUT ($o = 1$) or not ($o = 0$).

YAHOO FINANCE DATASET

The [Yahoo Finance dataset](#) is a collection of financial data on various assets traded in the financial markets. This dataset provides historical pricing information and range of financial metrics including cryptocurrencies.

In python, the Yahoo Finance dataset is accessed using the [yfinance](#) library. You will need to install it first using `pip install yfinance`.

Downloading Ethereum Data from Yahoo Finance: The line of code below uses the download function from the [yfinance](#) library to download historical price data of ETH against the USD currency pair. The data is downloaded for the time period from [January 1, 2018, to January 1, 2024](#).

```
1 eth_data = yf.download('ETH-USD', start='2018-01-01', end='2024-01-01')
```


YAHOO FINANCE DATASET

The dataset gives valuable information to explore Ethereum's market history. Each row represent a day. The fields (features) represent the ETH-USD price of the day including the highest, lowest and closing ETH-USD price for a day represented as high, low and close respectively.

Date	# Open	# High	# Low	# Close	# Adj Close	# Volume
2018-01-01	755.7570190429688	782.530029296875	742.0040283203125	772.6409912109375	772.6409912109375	2595760128
2018-01-02	772.3460083007812	914.8300170898438	772.3460083007812	884.4439697265625	884.4439697265625	5783349760
2018-01-03	886.0	974.4710083007812	868.4509887695312	962.719970703125	962.719970703125	5093159936
2018-01-04	961.7130126953125	1045.0799560546875	946.0859985351562	980.9219970703125	980.9219970703125	6502859776
2018-01-05	975.75	1075.3900146484375	956.3250122070312	997.719970703125	997.719970703125	6683149824
2018-01-06	995.1539916992188	1060.7099609375	994.6220092773438	1041.6800537109375	1041.6800537109375	4662219776

PREPROCESSING

The [preprocessing](#) prepare the Ethereum price data by performing two essential tasks. First, it extracts the [closing prices](#) of Ethereum, focusing specifically on this key metric for analysis. Then, utilizing the extracted data, using [min-max](#) scaling to normalize the closing prices, ensuring that they fall within the range of [\[0, 1\]](#).

```
1 def preprocessing(self):  
2     # Processing the data  
3     data = self.eth_data['Close'].values.reshape(-1, 1)  
4     self.scaler = MinMaxScaler(feature_range=(0,1))  
5     self.scaled_data = self.scaler.fit_transform(data)
```

Next, [TimeseriesGenerator](#) is used to construct generators for both training and testing datasets. Each generator produce sequences of input-output pairs, where both the input and output sequences are derived from the scaled Ethereum price data. The [length](#) parameter specifies the length of each input sequence, while the [batch_size](#) parameter controls the number of samples per batch.

MODEL

Layer (type)	Output Shape	Param
=====		
lstm (LSTM)	multiple	10400
lstm_1 (LSTM)	multiple	20200
dense (Dense)	multiple	51
=====		

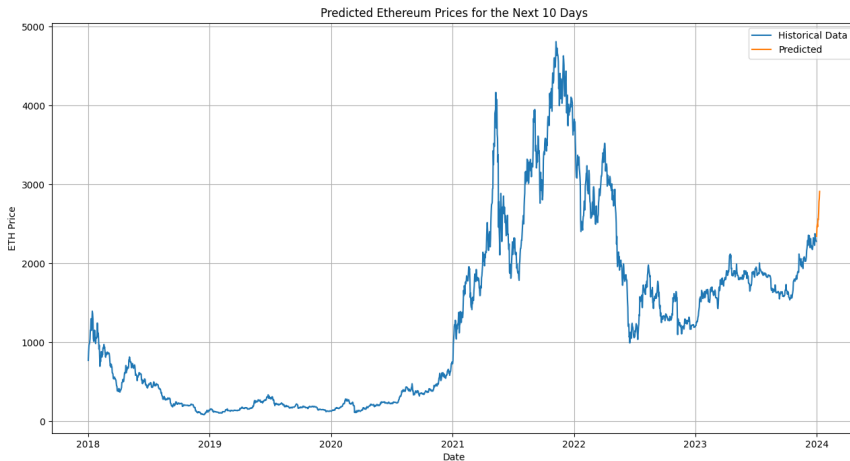
Total params: 30,651

Trainable params: 30,651

Non-trainable params: 0

For LSTM layers, the parameter count is determined by the number of units (50 in each LSTM layer) and the number of input features (1 in this case). The calculation involves multiple factors, including weights, biases, and gates within the LSTM cells. The Dense layer has 51 parameters, which includes the weight matrix of size (50, 1) and the bias vector of size (1,). The extra parameter accounts for the bias term.

ETH PRICE PREDICTION FOR NEXT 10 DAYS



RESOURCES

To download the source codes used in the previous slides, follow the link:

► [Download Source Codes](#)

Import the codes into your preferred development environment, such as Visual Studio Code (VS Code), to practice and explore further.

To learn programming in Python, follow my comprehensive 15-week Programming in Python course at:

► [Programming in Python](#)